



Deep learning with Pytorch

Mansura Habiba

Advisory Software Engineer

IBM Talent Management Solutions

Courtesy: [PyLadies Dublin](#)

Agenda

01

Introduction to
pytorch

02

Building blocks
of pytorch

03

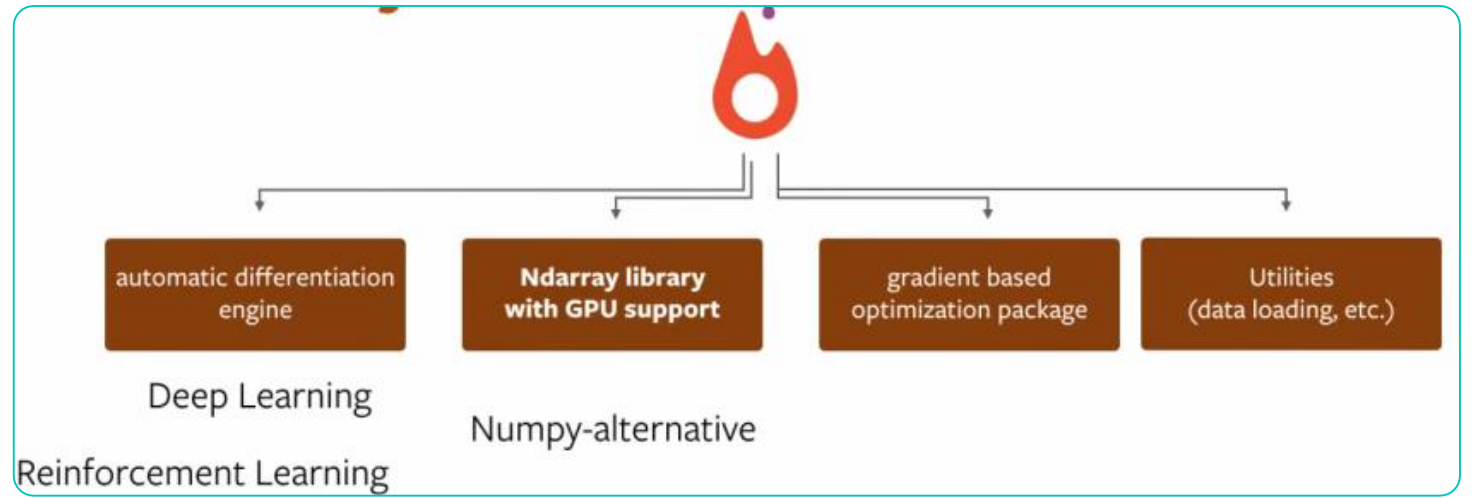
Deep learning
with pytorch

- Linear regression
- Sequence generation using NN

04

pytorch and
production

Introduction to Pytorch



- Lua torch
- C –libraries
- Automatic Differentiation Engine
- Nd Library with GPU
- Open-sourced by Facebook 2017

Installation of pytorch



Anaconda



Git repo



Docker Image




Helm Kubernetes

Setup Docker Container

Install Docker	<code>docker pull pytorch/pytorch</code>
Run the docker image	<code>docker run -it --name pytorch1 -v current_dir_including_ipynb_files:/workspace -p 5000:8888 -p 5001:6006 pytorch/pytorch</code>
Check the libraries	<code>pip freeze</code>
Install jupyter	<code>pip install jupyter</code>
Run Jupyter	<code>jupyter notebook --ip 0.0.0.0 --port 8888 --allow-root &</code>

Building blocks for deep learning in Pytorch



Brief introduction Tensor

't'
'e'
'n'
's'
'o'
'r'

tensor of dimensions [6]
(vector of dimension 6)

3	1	4	1
5	9	2	6
5	3	5	8
9	7	9	3
2	3	8	4
6	2	6	4

tensor of dimensions [6,4]
(matrix 6 by 4)

tensor of dimensions [4,4,2]

1
Rank:0

$\begin{bmatrix} 1 \\ 2 \end{bmatrix}$

Rank:1

$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

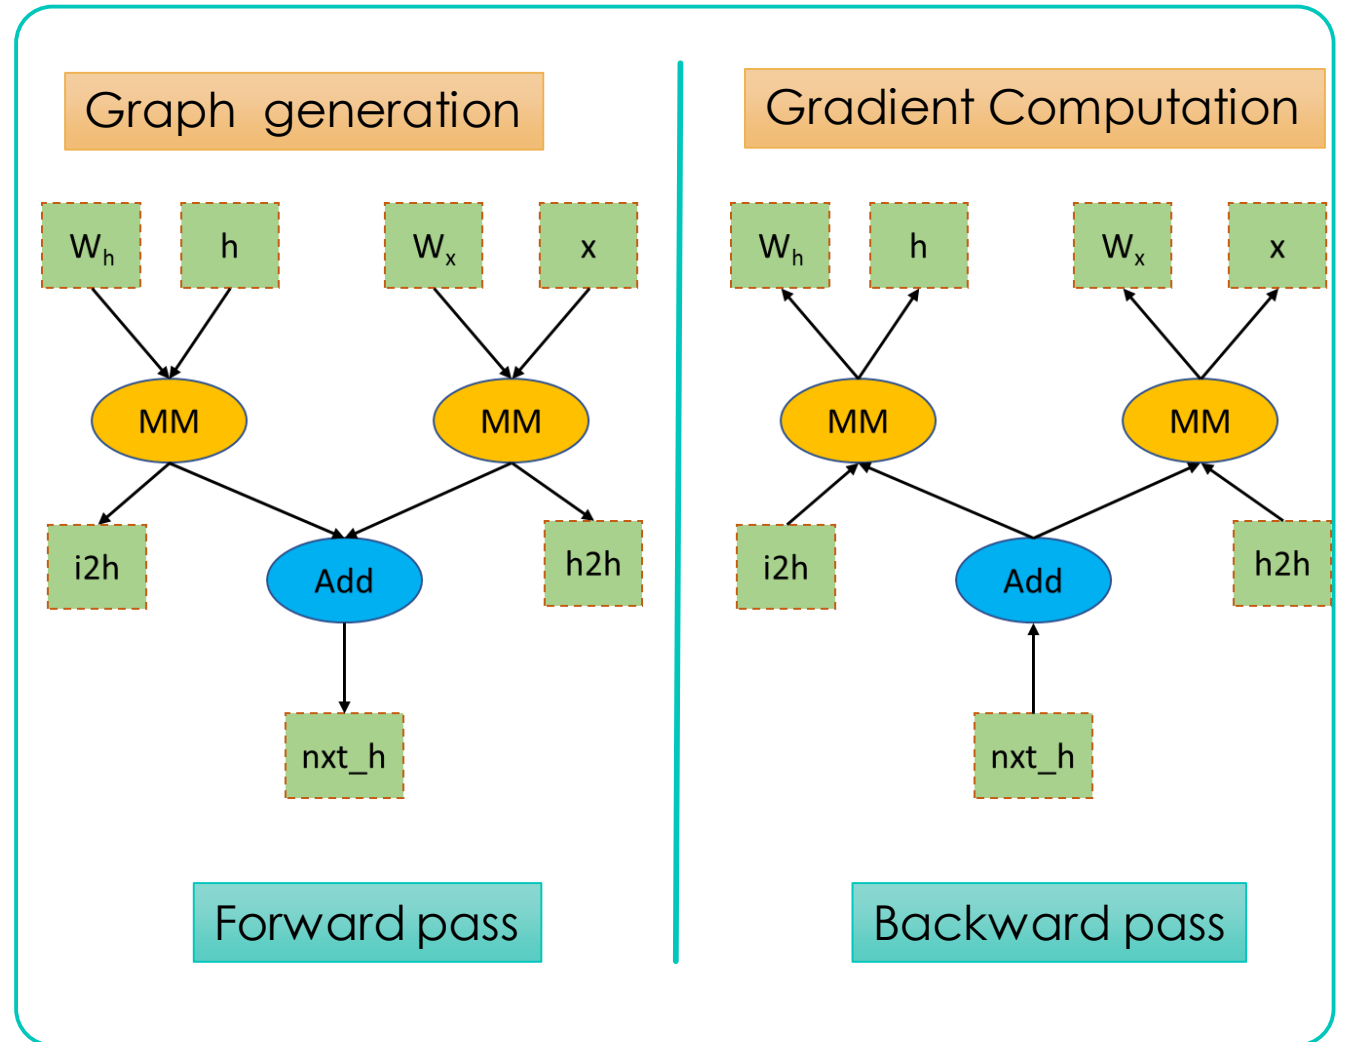
Rank:2

$\begin{bmatrix} 1 & 2 & 3 & 2 \\ 1 & 7 & 5 & 4 \end{bmatrix}$

Rank:3

Dynamic Computation Graph

- Why dynamic?
- `retain_graph = true`



Computational graph toolkit

```
1 import tensorflow as tf
2 import numpy as np
3
4 trX = np.linspace(-1, 1, 101)
5 trY = 2 * trX + np.random.randn(*trX.shape) * 0.33
6
7 X = tf.placeholder("float")
8 Y = tf.placeholder("float")
9
10 def model(X, w):
11     return tf.multiply(X, w)
12
13 w = tf.Variable(0.0, name="weights")
14 y_model = model(X, w)
15
16 cost = tf.square(Y - y_model)
17
18 train_op = tf.train.GradientDescentOptimizer(0.01).minimize(cost)
19
20 with tf.Session() as sess:
21     tf.global_variables_initializer().run()
22
23     for i in range(100):
24         for (x, y) in zip(trX, trY):
25             sess.run(train_op, feed_dict={X: x, Y: y})
26
27 print(sess.run(w))
```

Input / output placeholders


Declarative toolkit

```
1 import torch
2 from torch.autograd import Variable
3
4 trX = torch.linspace(-1, 1, 101)
5 trY = 2 * trX + torch.randn(*trX.size()) * 0.33
6
7 w = Variable(trX.new([0.0]), requires_grad=True)
8
9 for i in range(100):
10     for (x, y) in zip(trX, trY):
11         X = Variable(x)
12         Y = Variable(y)
13
14         y_model = X * w.expand_as(X)
15         cost = (Y - Y_model) ** 2
16         cost.backward(torch.ones(*cost.size()))
17
18         w.data = w.data + 0.01 * w.grad.data
19
20 print(w)
```

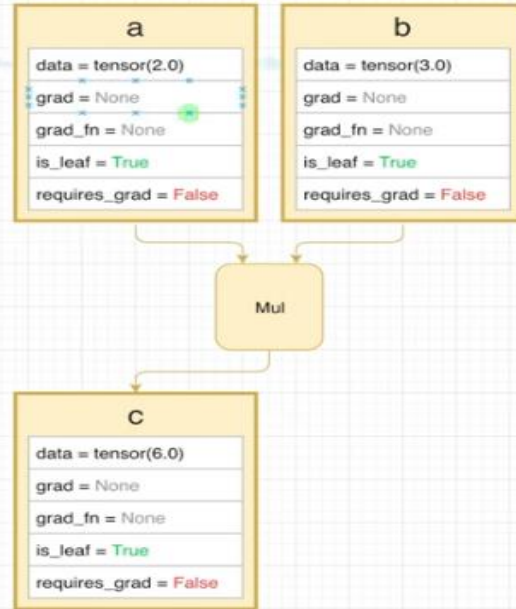
Imperative toolkit

Variable

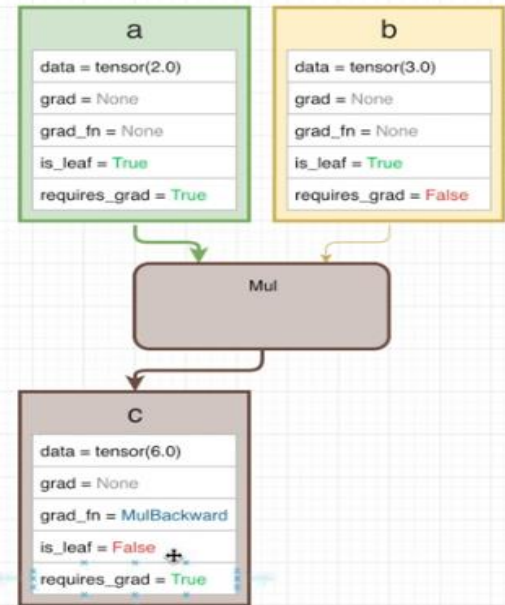
- Wrapper around a tensor
- It store history all the operations done on the tensor
- Main purposes are
 - Computation graph specification
 - Accumulation of gradients
 - Facilitate Back propagation, automatic differentiation
- `requires_grad = True`
- `torch.no_grad()`



```
a = torch.tensor(2.0)
b = torch.tensor(3.0)
```



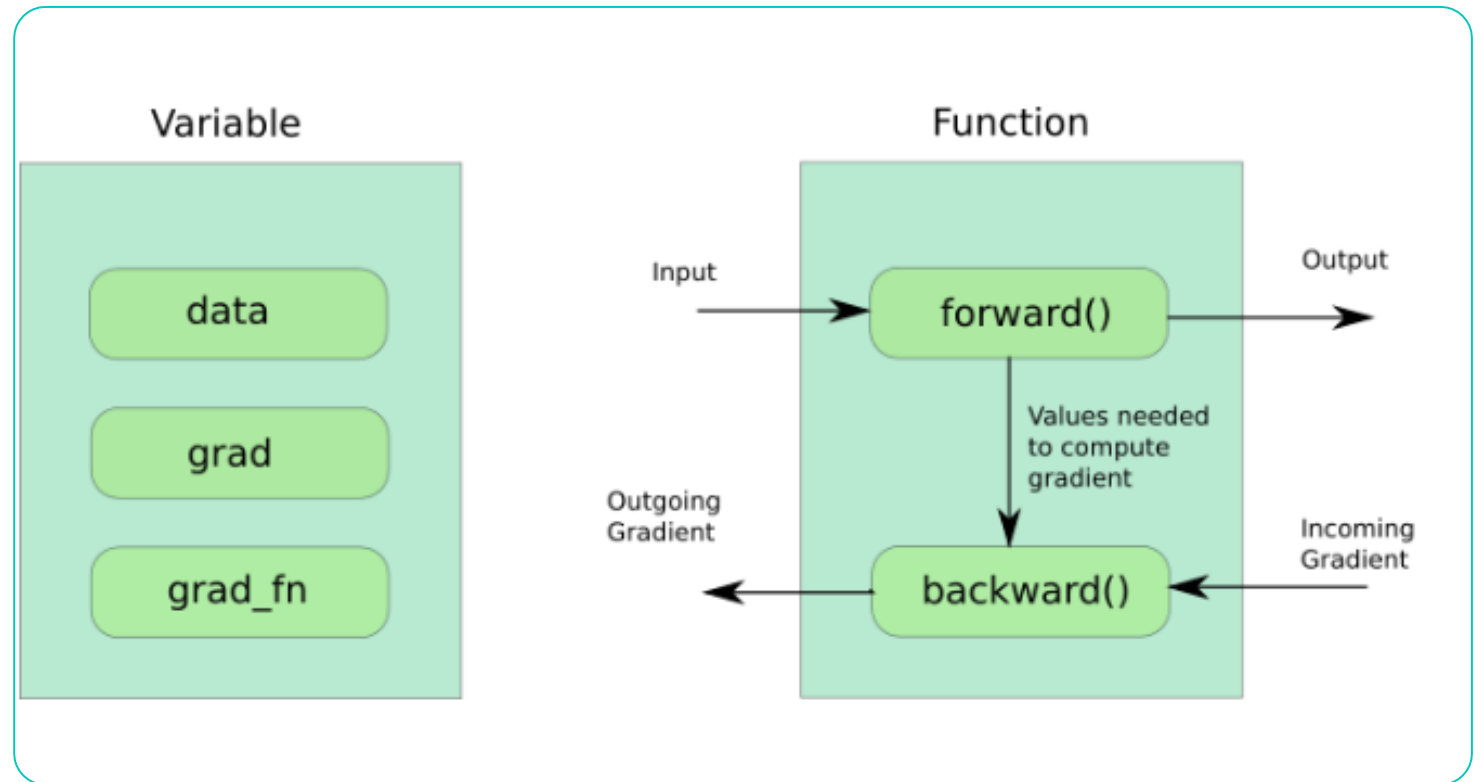
```
a = torch.tensor(2.0,
requires_grad=True)
b = torch.tensor(3.0)
```



requires_grad=False

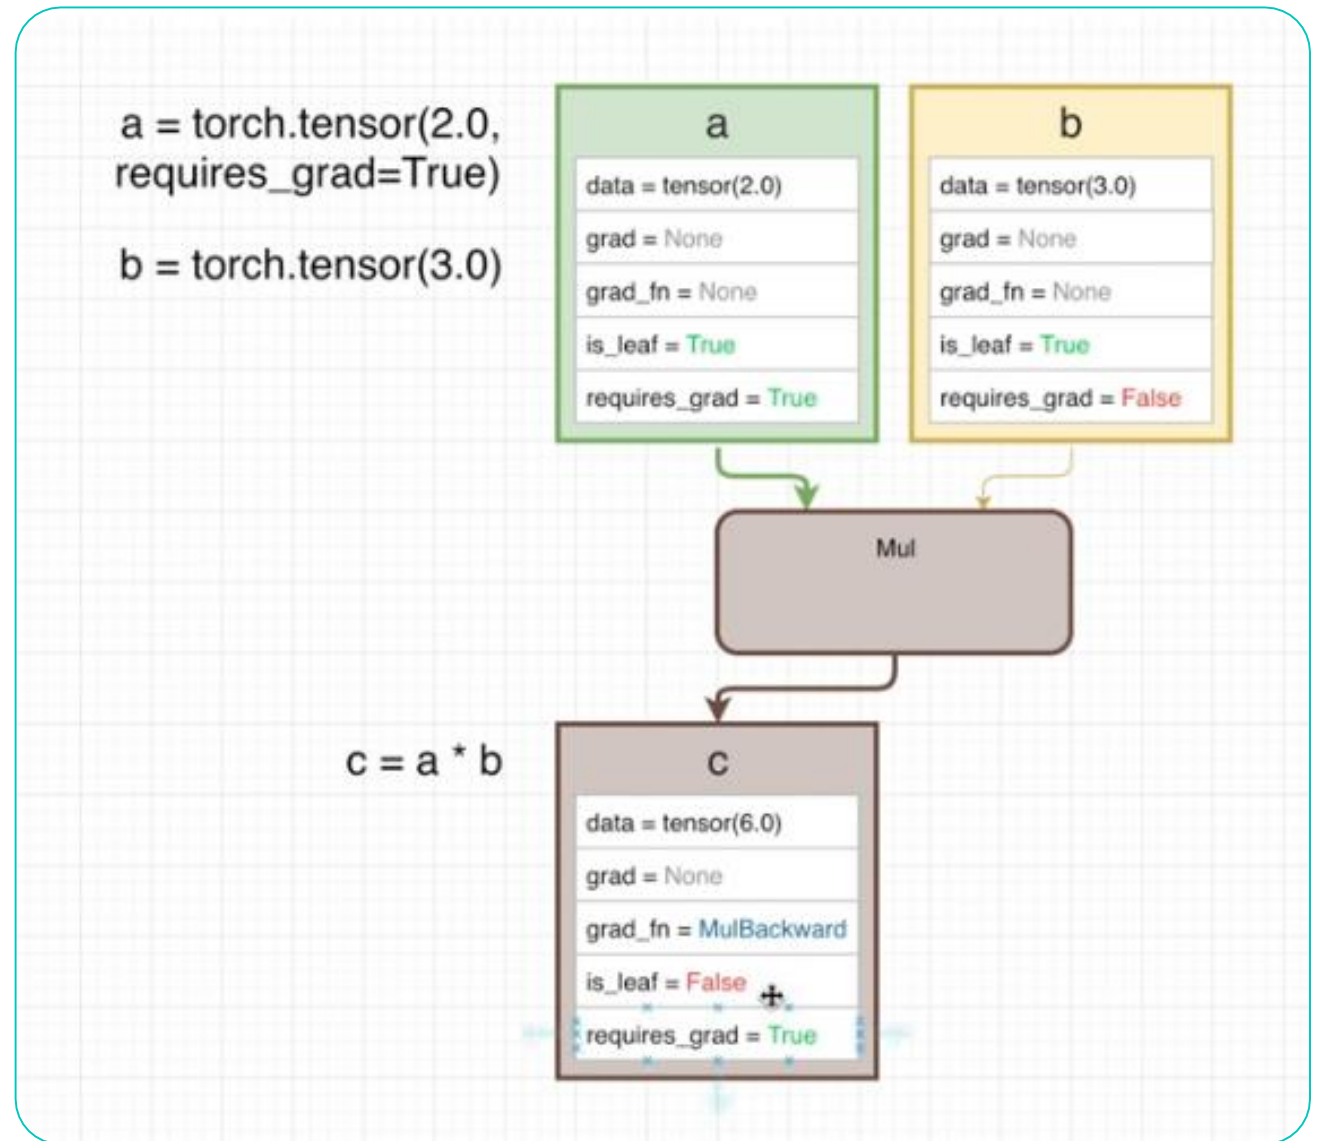
Function

- $\text{next_h} = \text{i2h} + \text{h2h}$
- $\text{grad_fn} = \text{AddBackWard}$



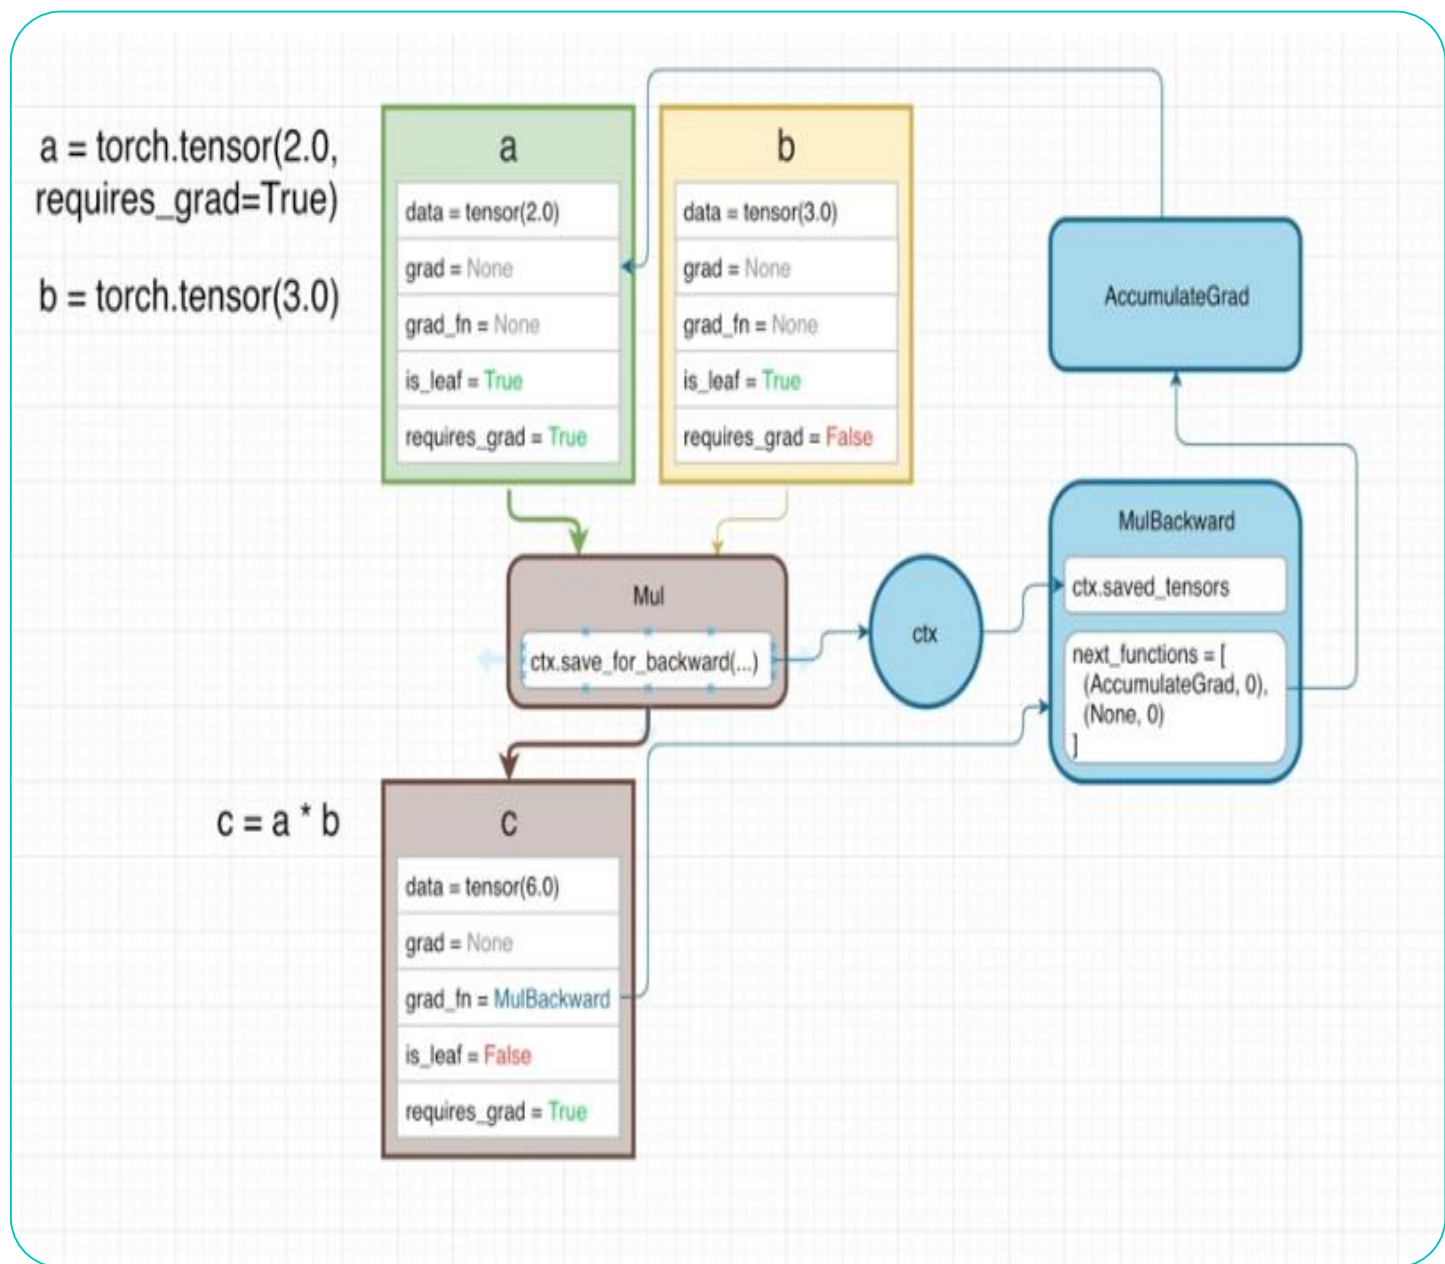
Autograd

- Forward pass
- Computation Graph generation



Autograd

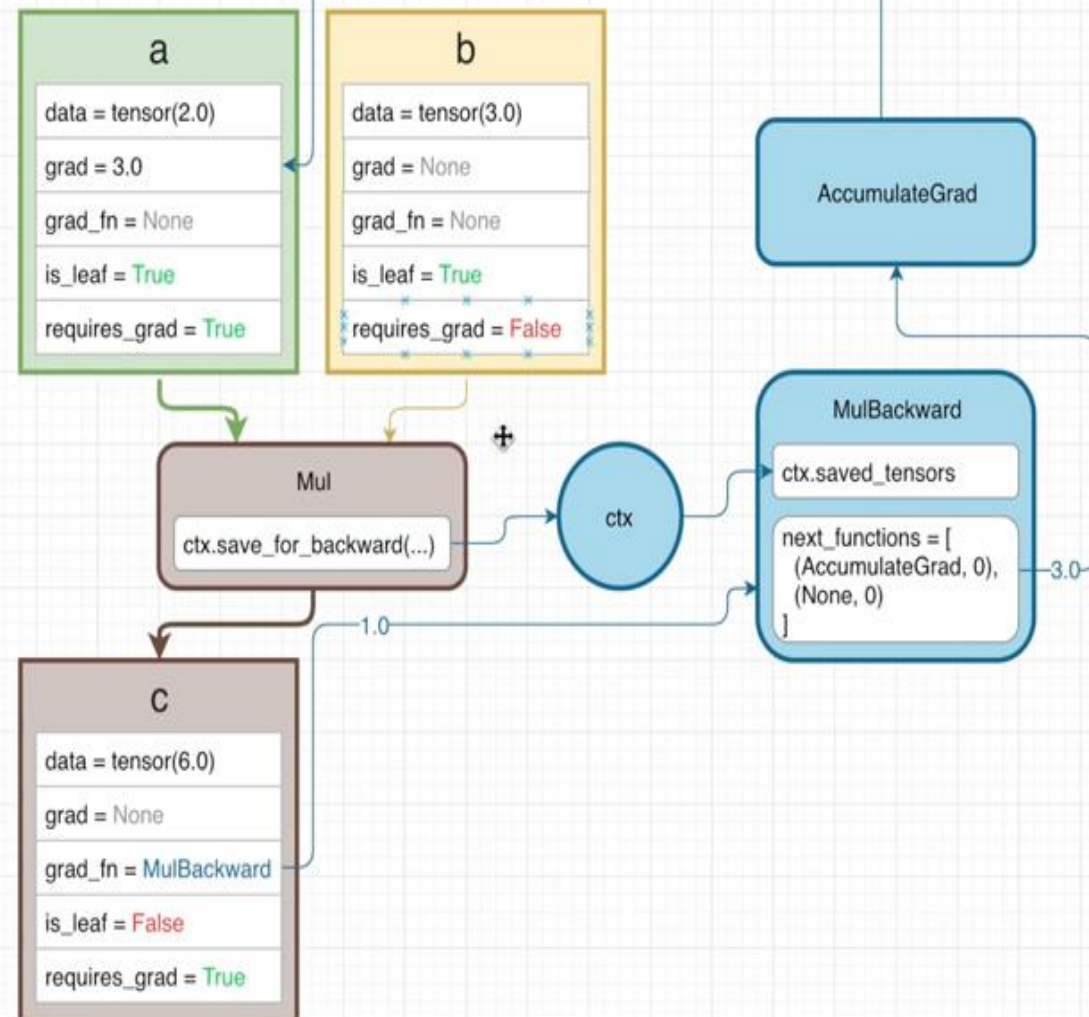
- Backward pass
- Computation gradient



Autograd

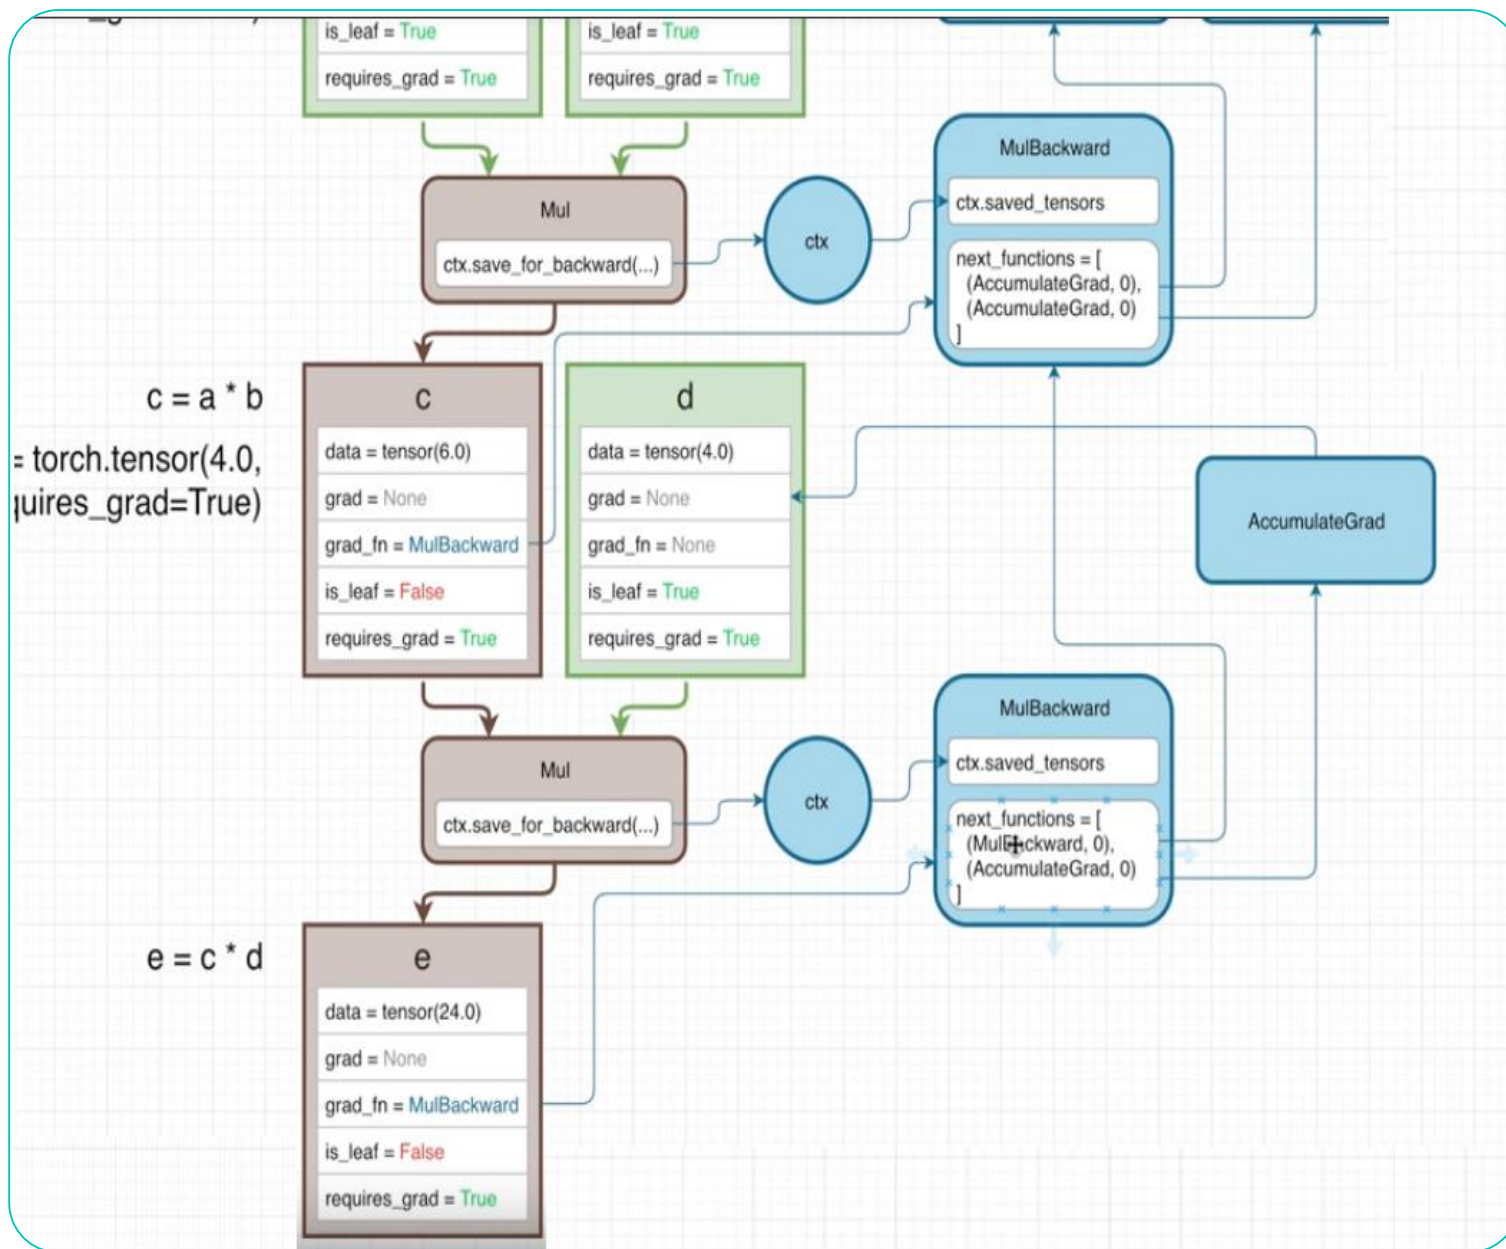
- Variable also stores the gradient of a scalar quantity (say, loss) with respect to the parameter it holds

```
a = torch.tensor(2.0,  
requires_grad=True)  
b = torch.tensor(3.0)
```



Autograd

- More deep tree

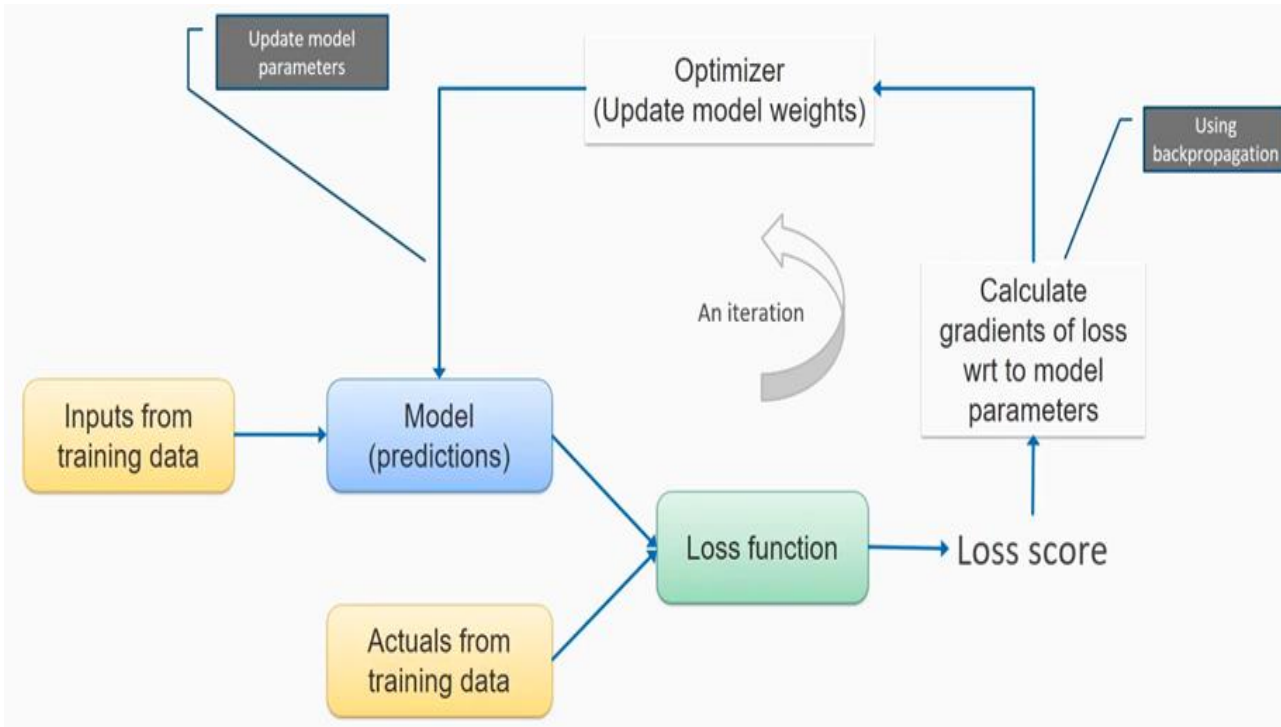


Pytorch and automatic differentiation – in a nutshell

```
from torch.autograd import Variable
x, prev_h = Variable(torch.randn(1, 10)), Variable(torch.randn(1, 20))
W_h, W_x = Variable(torch.randn(20, 20)), Variable(torch.randn(20, 10))

i2h = torch.matmul(W_x, x.t())
h2h = torch.matmul(W_h, prev_h.t())
(i2h + h2h).tanh().sum().backward()
```

- AD for pytorch written in C++
- Every intermediate result records only the subset of the computation graph that was relevant to their computation
- Support Invalidation and Aliasing



```
for input, target in dataset:  
    optimizer.zero_grad()  
    output = model(input)  
    loss = loss_fn(output, target)  
    loss.backward()  
    optimizer.step()
```

Role of Optimizer

SGD, Adagrad, RMSProp, LBFGS etc.

Why do we need to call zero_grad() in PyTorch?

- In PyTorch, we need to set the gradients to zero before starting to do backpropagation because PyTorch accumulates the gradients on subsequent backward passes.
- So, the default action is to accumulate the gradients on every `loss.backward()` call.

```
import torch
from torch.autograd import Variable
import torch.optim as optim

def linear_model(x, W, b):
    return torch.matmul(x, W) + b

data, targets = ...

W = Variable(torch.randn(4, 3), requires_grad=True)
b = Variable(torch.randn(3), requires_grad=True)

optimizer = optim.Adam([W, b])

for sample, target in zip(data, targets):
    # clear out the gradients of all Variables
    # in this optimizer (i.e. W, b)
    optimizer.zero_grad()
    output = linear_model(sample, W, b)
    loss = (output - target) ** 2
    loss.backward()
    optimizer.step()
```

Why Pytorch is fast



Avoid GIL by JIT



Use Cuda/ GPU

Main Modules of Pytorch

Neural Network Model

`torch.nn`


Autograd

`torch.autograd`

Optimizer

`torch.optim`

Linear regression



```
import torch
```

```
w=torch.tensor(2.0,requires_grad=True)  
b=torch.tensor(-1.0,requires_grad=True)
```

```
def forward(x):  
    y=w*x+b  
    return y
```

```
x=torch.tensor([[1.0]])
```

```
yhat=forward(x)
```

```
yhat: tensor([[1.0]])
```

LR in General way

○ $b = -1, w = 2$

○ $y = -1 + 2x$

○ $x = 1$

○ $y = -1 + 2(1)$

○ $y = 1$


```
import torch
```

```
w=torch.tensor(2.0,requires_grad=True)  
b=torch.tensor(-1.0,requires_grad=True)
```

```
def forward(x):  
    y=w*x+b  
    return y
```

```
x=torch.tensor([[1.0]])
```

```
yhat=forward(x)
```

```
yhat: tensor([[1.0]])
```

```
from torch.nn import Linear as Linear
```

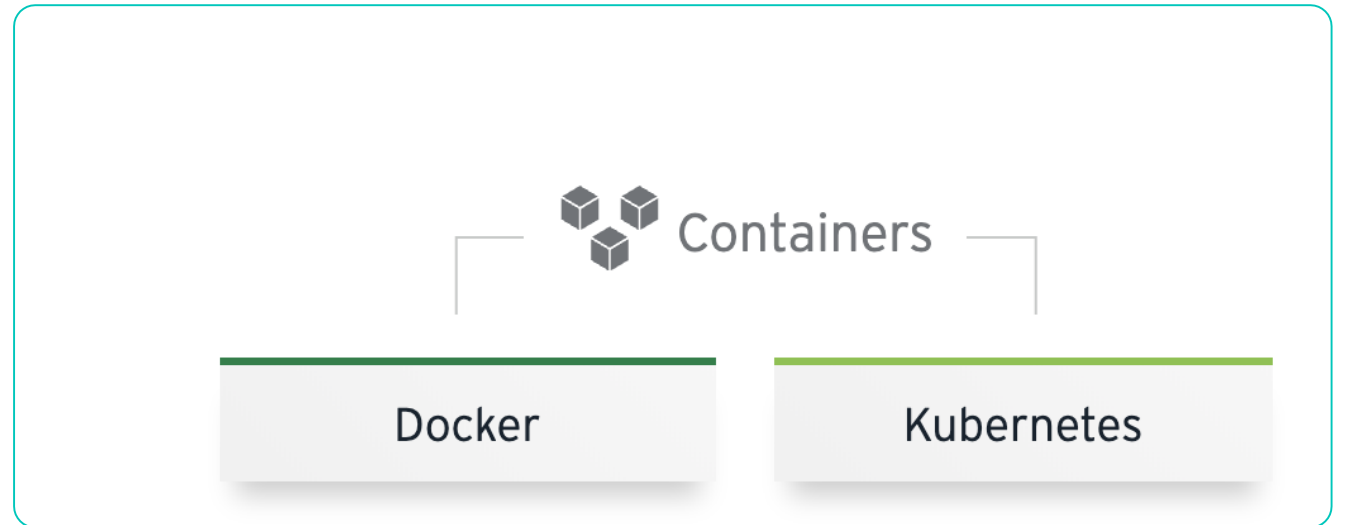
```
model=Linear(in_features=1,out_features=1)
```

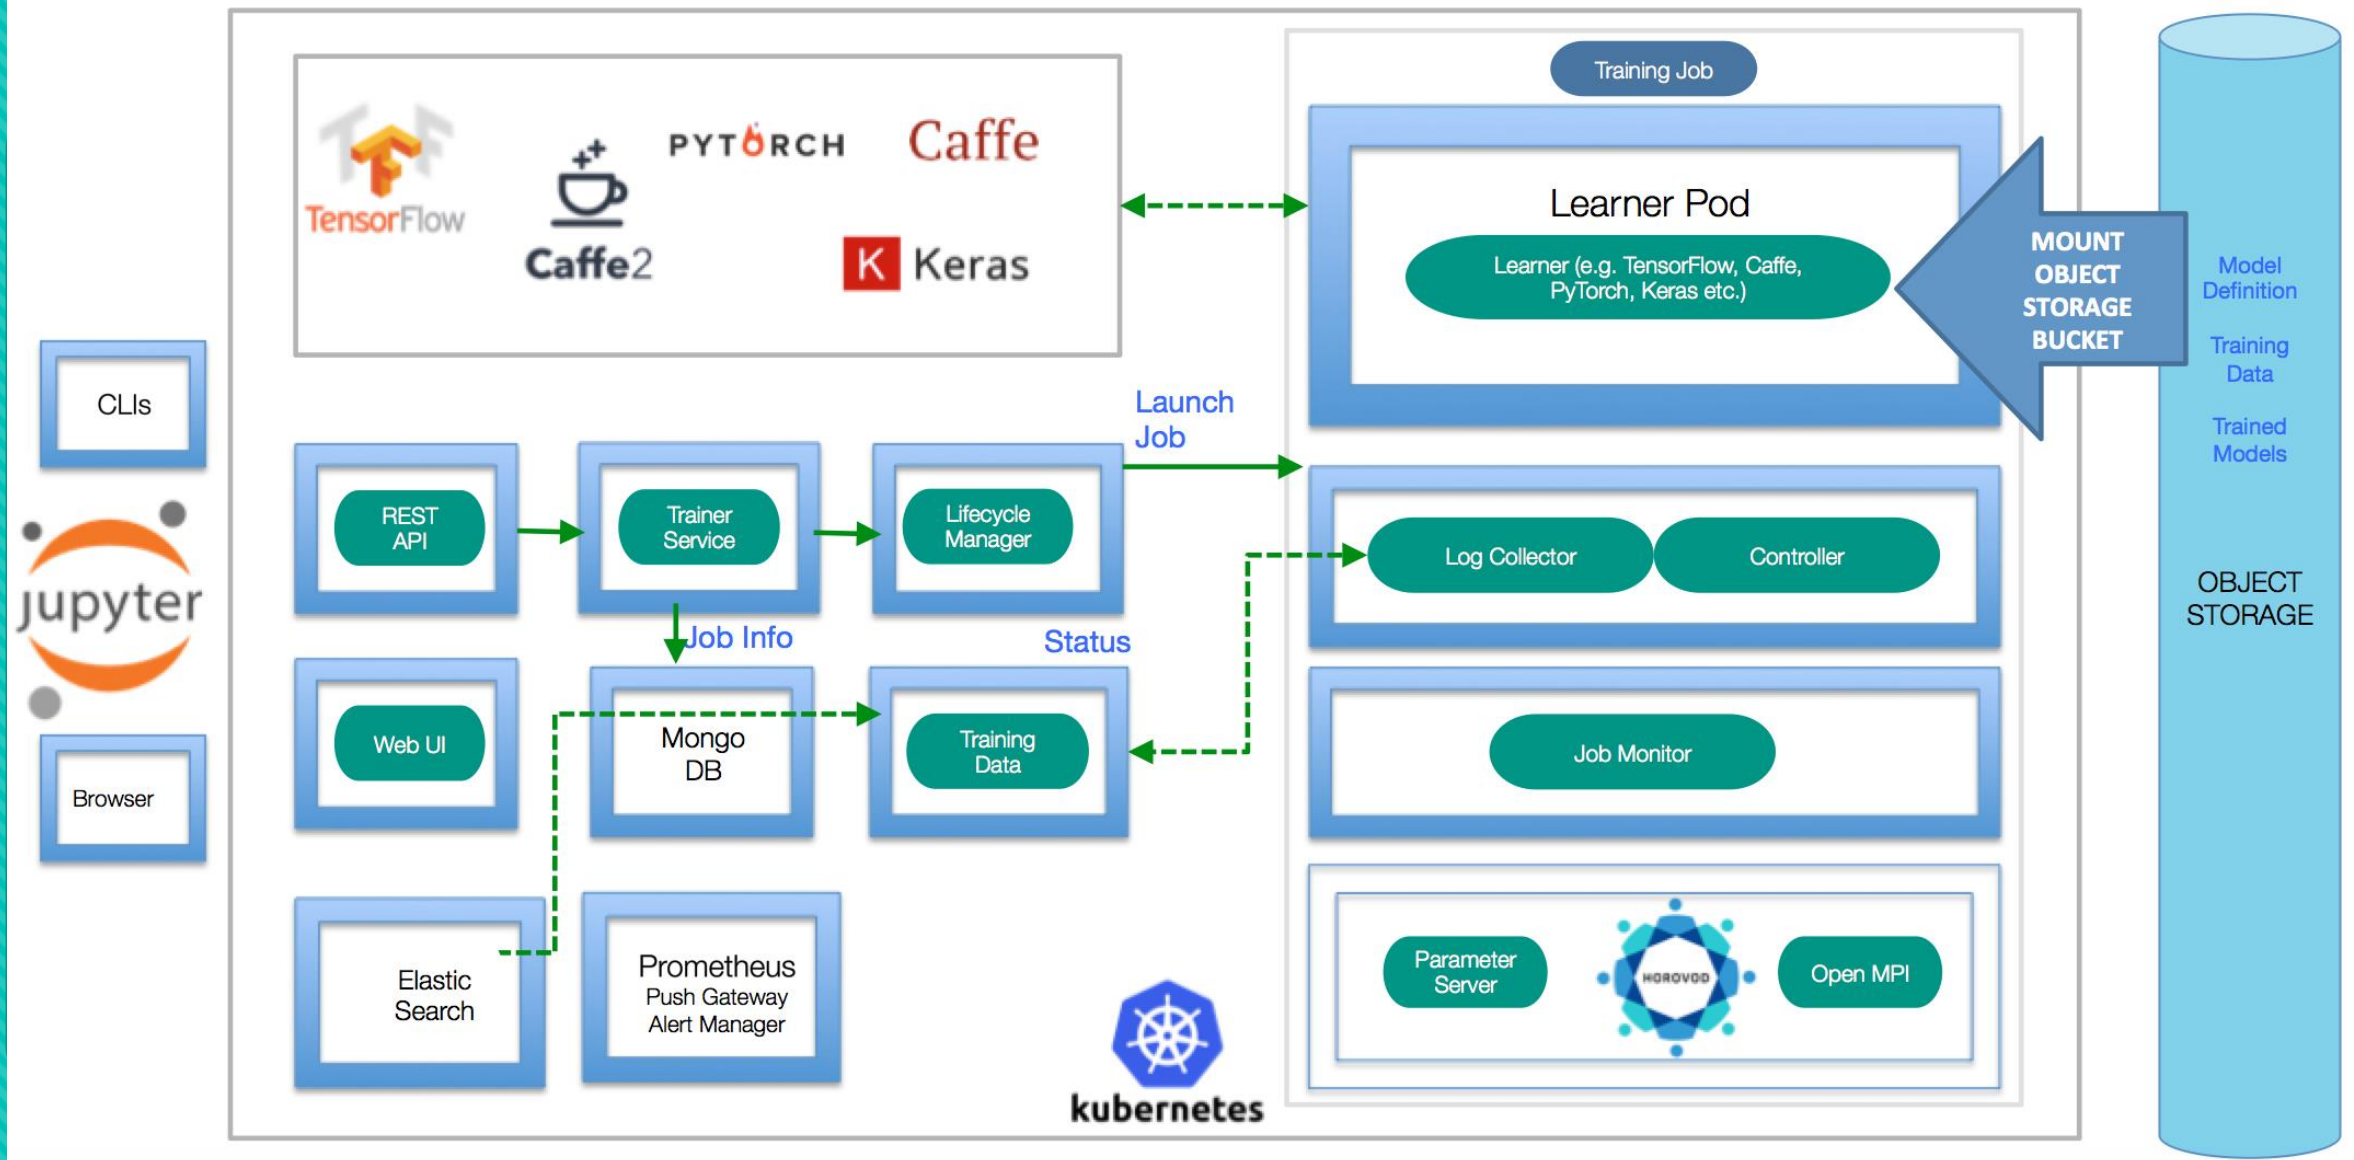
```
yhat=model(x)
```

LR in Pytorch way

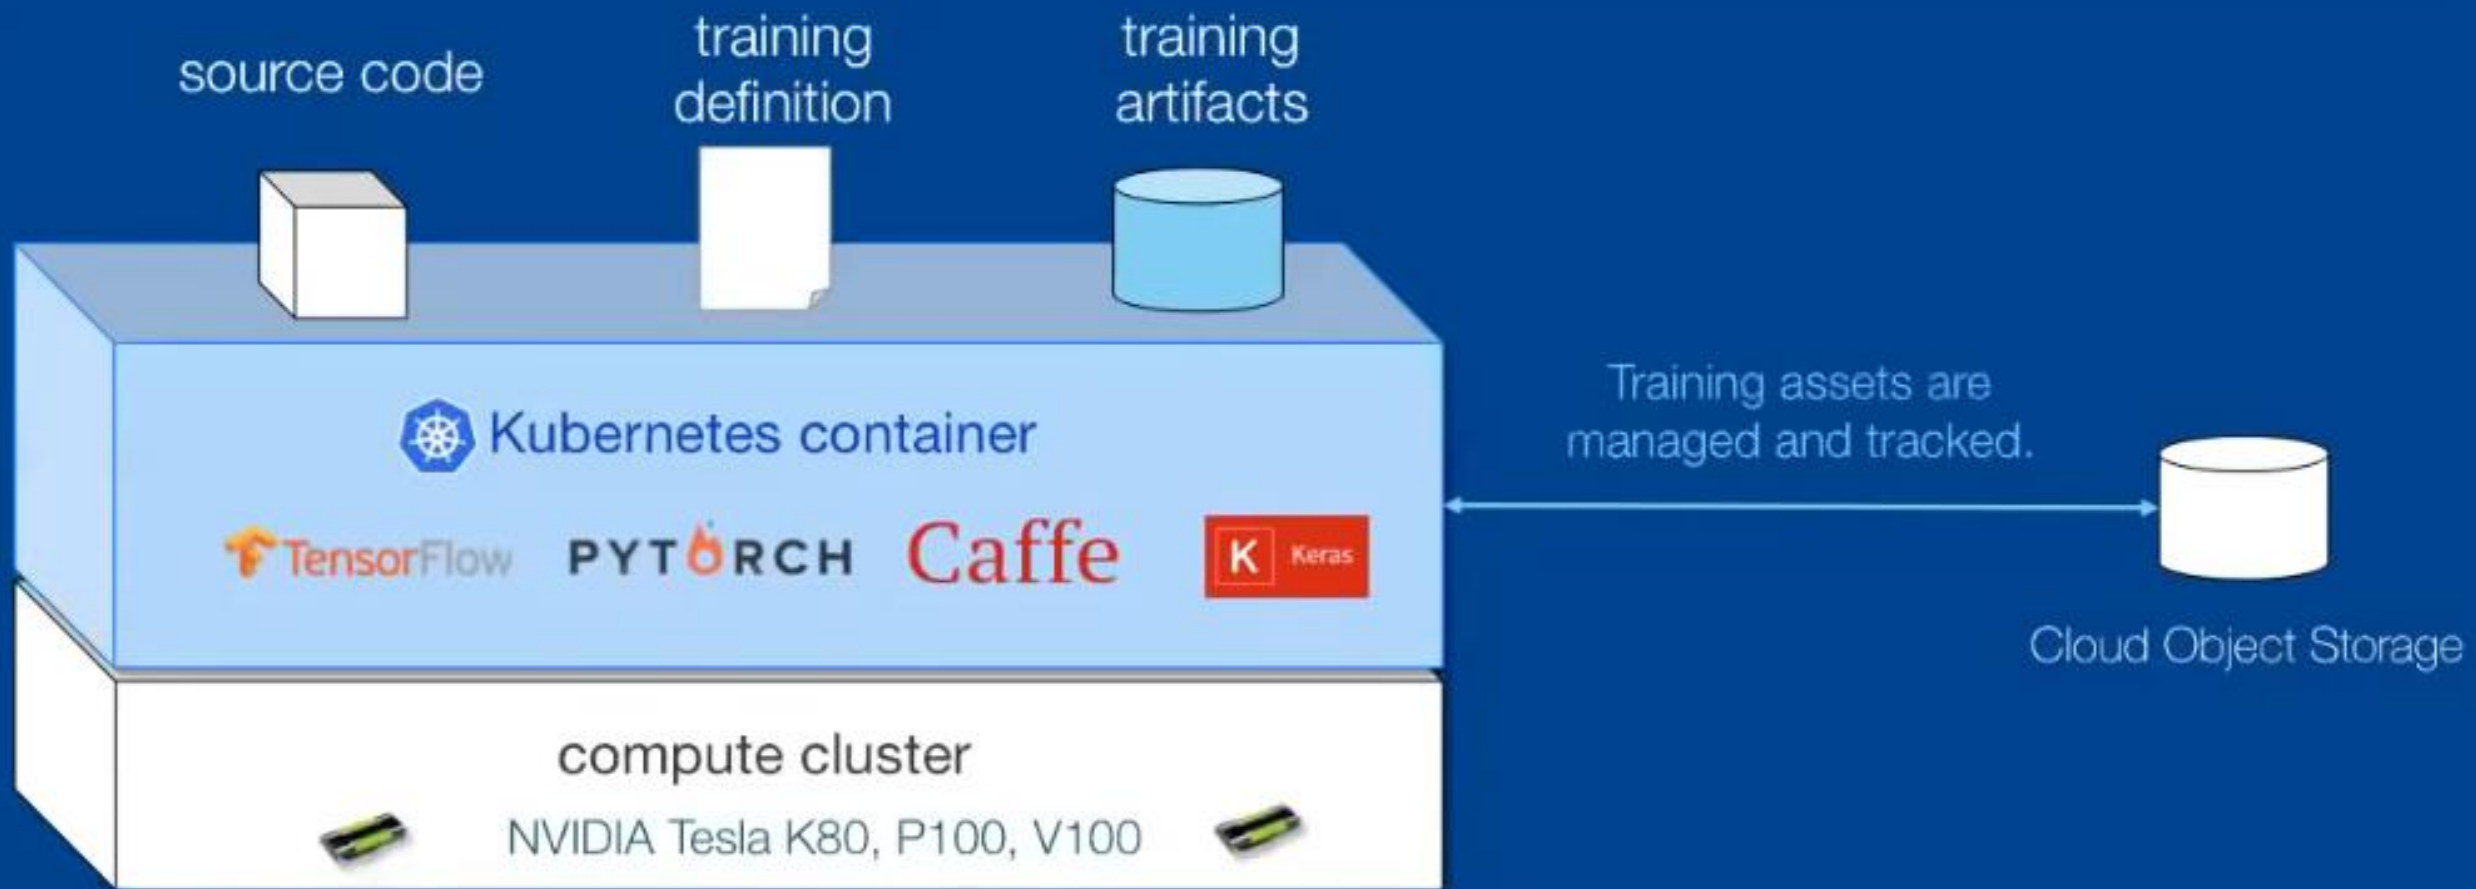
Choices for Production

- Save the model `torch.save`
- Load the model `torch.load`






Deep Learning ecosystem for production : IBM FfDL



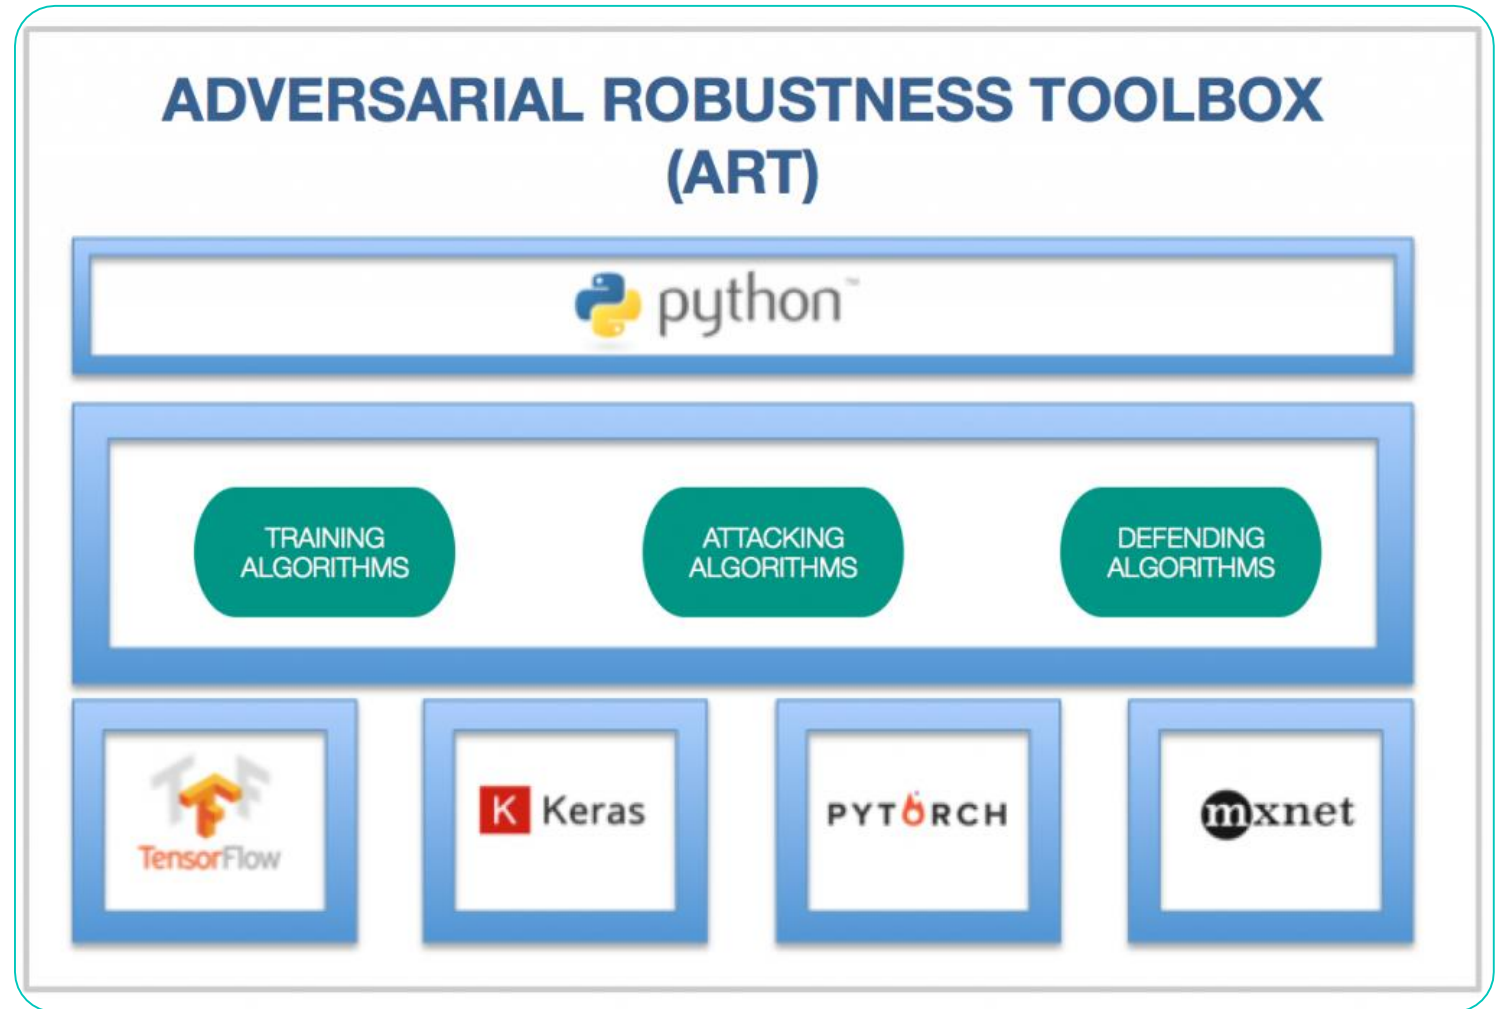
IBM FfDL



Training Model using FfDL

Popular Library of FfDI

ART –
(<https://github.com/IBM/adversarial-robustness-toolbox>)



Start With IBM:FfDL

01

Getting Stared

02

Sample Model

03

Pytorch Model

Reference

- Paul O'Grady - An introduction to PyTorch & Autograd

Learning Materials



Official tutorials:
<https://pytorch.org/tutorials/>



Examples:
<https://github.com/pytorch/examples>



Course : [Deep Learning with Python and PyTorch](#)



Playlist :
<http://deeplizard.com/learn/video/v5cngxo4mlg>



Podcast :
<https://podtail.com/en/podcast/this-week-in-machine-learning-ai-podcast/pytorch-fast-differentiable-dynamic-graphs-in-pyth/>



Thank you

Partial Derivative

$$f = uv + u^2$$

$$f(u = 1, v = 2) = uv + u^2$$

$$1(2) + 1^2 = 3$$

$$\frac{\partial f(1,2)}{\partial u}, \frac{\partial f(1,2)}{\partial v}$$

$$\frac{\partial f(u, v)}{\partial u} = v + 2u$$

$$\frac{\partial f(u = 1, v = 2)}{\partial u} = 2 + 2(1)$$

Derivative of u

$$f(u = 1, v = 2) = uv + u^2$$

$$1(2) + 1^2 = 3$$

$$\frac{\partial f(1,2)}{\partial u}, \frac{\partial f(1,2)}{\partial v}$$

$$\frac{\partial f(u, v)}{\partial v} = u$$

$$\frac{\partial f(u = 1, v = 2)}{\partial v} = 1$$

$$= 1$$

Derivative of v